



1a Human-authored story



1b LLM preparation



2 Agent investigation

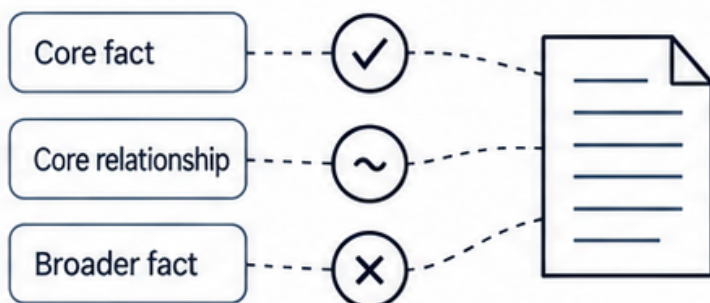


3 Monitor reconstruction



5 configurations • 250 reports

4 Rubric-based grading



2 LLM graders + 1 LLM reviewer

Correct Partial Incorrect Not reported Unresolved

5 Numerical scores

Example aggregate: **Astra high**

45.6%

Story reconstruction

Fully recovered rubric items

78.0%

Culprit accuracy

Correctly identified reports

Scenario +
role packs

Interaction logs

Monitor
reports

Reference rubric