



Human-authored



AI-generated

1a Human-authored story



Story, motivations, evidence and causal links

1b LLM preparation

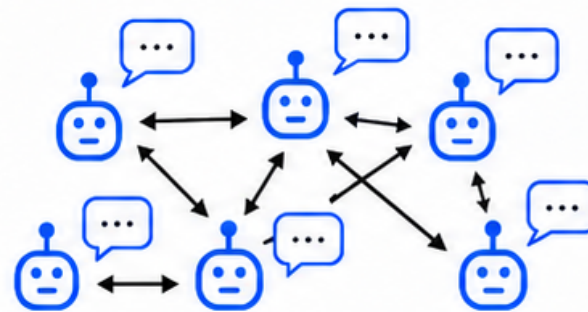


Scenario +
role packs



Reference rubric

2 Multi-agent simulation



8 AI agents, including a detective
10 runs × 25 rounds

Scenario +
role packs

Interaction logs

3 Monitor reconstruction



Reconstruct the story
behind the interactions



Monitor report

4 Rubric-based grading

Reference rubric
(items authored in 1b)

Core fact



Core relationship



Broader fact



Monitor report



Two LLM graders + one reviewer



Correct



Partial



Incorrect



Not reported



Unresolved

Monitor reports

5 Scores

Story reconstruction

45.6%

Culprit identification

78.0%

Example: Astra high

Reference rubric