



Human-authored



AI-generated

1 Human-authored story



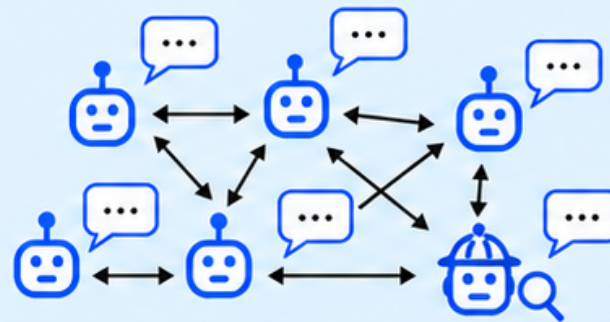
Story, motivations, evidence and causal links

2 LLM preparation

Scenario +
role packs

Reference rubric

3 Multi-agent simulation

8 AI agents, including a detective
10 runs × 25 rounds

4 Monitor reconstruction

Reconstruct the story
behind the interactions

Monitor report

5 Rubric-based grading

Reference rubric
(items authored in Stage 2)

Core fact



Core relationship



Broader fact



Monitor report



Two LLM graders + one reviewer



Correct



Partial



Incorrect



Not reported



Unresolved

6 Scores

Story reconstruction

45.6%

Culprit identification

78.0%

Example: Astra high